

Handling Energy Consumption Data with Pandas: Missing Values and Data Preprocessing

Objectives:

1. Handling Missing Values: We will demonstrate how to handle missing values in energy consumption data by removing rows/columns, imputing values using mean/median, applying forward/backward filling, and flagging missing data.
2. Data Preprocessing: We will normalize/standardize the data, encode categorical variables, and implement feature engineering for further analysis.

Step 1: Import Pandas and Create a Dataset with Missing Values

We'll first import the necessary libraries and create a dataset that includes some missing values (NaN).

```
In [1]: import pandas as pd
import numpy as np

# Sample data with missing values
data = {
    "Energy Source": ["Solar", "Wind", "Hydropower", "Geothermal", "Biomass"],
    "Energy Consumption (MWh)": [1200, np.nan, 2900, np.nan, 2500, 3200],
    "Cost (Million $)": [200, 400, np.nan, 150, 250, np.nan]
}

# Create a DataFrame
energy_df = pd.DataFrame(data)

print("Original Energy Data with Missing Values:")
print(energy_df)
```

Original Energy Data with Missing Values:

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)
0	Solar	1200.0	200.0
1	Wind	NaN	400.0
2	Hydropower	2900.0	NaN
3	Geothermal	NaN	150.0
4	Biomass	2500.0	250.0
5	Nuclear	3200.0	NaN

We created a Pandas DataFrame `energy_df` representing energy sources and their consumption and costs.

The dataset includes some NaN (missing) values, which we will handle in the next steps.

1. Handling Missing Values

1.1. Remove Rows with Missing Values

We can remove rows that contain any missing values using `dropna()`.

```
In [2]: # Remove rows with any missing values
cleaned_df = energy_df.dropna()

print("\nData After Removing Rows with Missing Values:")
print(cleaned_df)
```

Data After Removing Rows with Missing Values:

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)
0	Solar	1200.0	200.0
4	Biomass	2500.0	250.0

The above code snippet removes rows where any column has missing data. This method is straightforward but may result in losing a significant amount of data.

1.2. Impute Missing Values with the Mean

Instead of removing rows, we can impute missing values by filling them with the mean value of the column.

```
In [3]: # Impute missing values in 'Energy Consumption (MWh)' with the mean
energy_df["Energy Consumption (MWh)"].fillna(energy_df["Energy Consumption (MWh)"].mean(), inplace=True)

# Impute missing values in 'Cost (Million $)' with the mean
energy_df["Cost (Million $)"].fillna(energy_df["Cost (Million $)"].mean(), inplace=True)

print("\nData After Imputing Missing Values with Mean:")
print(energy_df)
```

Data After Imputing Missing Values with Mean:

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)
0	Solar	1200.0	200.0
1	Wind	2450.0	400.0
2	Hydropower	2900.0	250.0
3	Geothermal	2450.0	150.0
4	Biomass	2500.0	250.0
5	Nuclear	3200.0	250.0

We used the mean imputation method to fill missing values in both the Energy Consumption (MWh) and Cost (Million \$) columns, ensuring that we retain the dataset while handling missing values.

1.3. Forward/Backward Filling

Another method is forward filling, where missing values are replaced by the previous valid entry.

```
In [4]: # Forward fill missing values
forward_filled_df = energy_df.fillna(method="ffill")

print("\nData After Forward Filling:")
print(forward_filled_df)
```

Data After Forward Filling:

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)
0	Solar	1200.0	200.0
1	Wind	2450.0	400.0
2	Hydropower	2900.0	250.0
3	Geothermal	2450.0	150.0
4	Biomass	2500.0	250.0
5	Nuclear	3200.0	250.0

Forward filling (ffill) replaces missing values with the previous non-missing value in the column, which is useful when data is time-series-based.

1.4. Flag Missing Values

We can also create a separate column to flag missing values before imputation.

```
In [5]: # Create a flag column indicating missing values in 'Energy Consumption (MWh)'
energy_df["Missing Consumption"] = energy_df["Energy Consumption (MWh)"].isna()

print("\nData with Missing Values Flagged:")
print(energy_df)
```

Data with Missing Values Flagged:

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)	\
0	Solar	1200.0	200.0	
1	Wind	2450.0	400.0	
2	Hydropower	2900.0	250.0	
3	Geothermal	2450.0	150.0	
4	Biomass	2500.0	250.0	
5	Nuclear	3200.0	250.0	

	Missing Consumption
0	0
1	0
2	0
3	0
4	0
5	0

The Missing Consumption column flags missing values with 1 (missing) or 0 (not missing), which helps track imputed values.

2. Data Preprocessing

2.1. Normalization (Min-Max Scaling)

We will scale the data to a range between 0 and 1 using Min-Max Scaling.

```
In [6]: from sklearn.preprocessing import MinMaxScaler

# Normalize the 'Energy Consumption (MWh)' and 'Cost (Million $)'
scaler = MinMaxScaler()
energy_df[["Energy Consumption (MWh)", "Cost (Million $)"]] = scaler.fit_transform(
    energy_df[["Energy Consumption (MWh)", "Cost (Million $)"]]
)

print("\nData After Normalization (Min-Max Scaling):")
print(energy_df)
```

Data After Normalization (Min-Max Scaling):

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)
0	Solar	0.000	0.2
1	Wind	0.625	1.0
2	Hydropower	0.850	0.4
3	Geothermal	0.625	0.0
4	Biomass	0.650	0.4
5	Nuclear	1.000	0.4

	Missing Consumption
0	0
1	0
2	0
3	0
4	0
5	0

Min-Max Scaling normalizes the energy consumption and cost values, scaling them to a range between 0 and 1. This is useful when comparing features with different units or magnitudes.

2.2. Standardization (Z-score Scaling)

Alternatively, we can use standardization, which centers the data around a mean of 0 with a standard deviation of 1.

```
In [7]: from sklearn.preprocessing import StandardScaler

# Standardize the 'Energy Consumption (MWh)' and 'Cost (Million $)'
scaler = StandardScaler()
energy_df[["Energy Consumption (MWh)", "Cost (Million $)"]] = scaler.fit_transform(
    energy_df[["Energy Consumption (MWh)", "Cost (Million $)"]]
)

print("\nData After Standardization (Z-score Scaling):")
print(energy_df)
```

Data After Standardization (Z-score Scaling):

	Energy Source	Energy Consumption (MWh)	Cost (Million \$)	\
0	Solar	-2.005893e+00	-6.546537e-01	
1	Wind	3.563181e-16	1.963961e+00	
2	Hydropower	7.221213e-01	1.817029e-16	
3	Geothermal	3.563181e-16	-1.309307e+00	
4	Biomass	8.023570e-02	1.817029e-16	
5	Nuclear	1.203536e+00	1.817029e-16	

	Missing Consumption
0	0
1	0
2	0
3	0
4	0
5	0

Z-score scaling standardizes the values, making the mean 0 and standard deviation 1, which is useful when dealing with normally distributed data.

2.3. Encoding Categorical Variables

We'll convert the categorical column Energy Source into numeric format using one-hot encoding.

```
In [8]: # One-hot encode the 'Energy Source' column
energy_encoded_df = pd.get_dummies(energy_df, columns=["Energy Source"])

print("\nData After One-Hot Encoding Categorical Variables:")
print(energy_encoded_df)
```

Data After One-Hot Encoding Categorical Variables:

	Energy Consumption (MWh)	Cost (Million \$)	Missing Consumption \
0	-2.005893e+00	-6.546537e-01	0
1	3.563181e-16	1.963961e+00	0
2	7.221213e-01	1.817029e-16	0
3	3.563181e-16	-1.309307e+00	0
4	8.023570e-02	1.817029e-16	0
5	1.203536e+00	1.817029e-16	0

	Energy Source_Biomass	Energy Source_Geothermal	Energy Source_Hydropower \
0	0	0	
0			
1	0	0	
0			
2	0	0	
1			
3	0	1	
0			
4	1	0	
0			
5	0	0	
0			

	Energy Source_Nuclear	Energy Source_Solar	Energy Source_Wind
0	0	1	0
1	0	0	1
2	0	0	0
3	0	0	0
4	0	0	0
5	1	0	0

One-hot encoding converts the Energy Source column into multiple binary columns, each representing the presence (1) or absence (0) of a specific energy source.

2.4. Feature Engineering

We can create a new feature that represents the ratio of energy consumption to cost.

```
In [9]: # Create a new feature: Energy Consumption per Million $
energy_encoded_df["Consumption per $Million"] = energy_encoded_df["Energy Co

print("\nData with New Feature (Consumption per $Million):")
print(energy_encoded_df)
```

Data with New Feature (Consumption per \$Million):

	Energy Consumption (MWh)	Cost (Million \$)	Missing Consumption \
0	-2.005893e+00	-6.546537e-01	0
1	3.563181e-16	1.963961e+00	0
2	7.221213e-01	1.817029e-16	0
3	3.563181e-16	-1.309307e+00	0
4	8.023570e-02	1.817029e-16	0
5	1.203536e+00	1.817029e-16	0

	Energy Source_Biomass	Energy Source_Geothermal	Energy Source_Hydropow
0	0	0	
1	0	0	
2	0	0	
3	0	1	
4	1	0	
5	0	0	

	Energy Source_Nuclear	Energy Source_Solar	Energy Source_Wind \
0	0	1	0
1	0	0	1
2	0	0	0
3	0	0	0
4	0	0	0
5	1	0	0

	Consumption per \$Million
0	3.064052e+00
1	1.814283e-16
2	3.974187e+15
3	-2.721424e-16
4	4.415764e+14
5	6.623646e+15

This new feature, Consumption per \$Million, calculates how much energy is produced per million dollars spent, providing insight into the efficiency of energy sources.

Conclusion

In this lab assignment, we handled missing values in energy consumption data by:

Removing rows with missing values,
Imputing missing values with the mean,
Using forward filling, and
Flagging missing values.

We then applied data preprocessing techniques such as normalization, standardization, encoding categorical variables, and feature engineering to enhance the dataset. These